

Student Performance Analysis - Classification based Project

```
In [47]: import pandas as pd
import seaborn as sns
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
```

```
In [48]: # Importing the dataset
```

```
In [83]: data = pd.read_csv("D:/HelloTech Softwares - Data Science Intern Projects/Student Performance dataset -
```

```
In [ ]: # Display first 10 rows
```

```
In [19]: data.head(10)
```

```
Out[19]:
```

	StudentID	Age	Gender	Ethnicity	ParentalEducation	StudyTimeWeekly	Absences	Tutoring	ParentalSupport	Extracurricular
0	1001	17	1	0	2	19.833723	7	1	2	0
1	1002	18	0	0	1	15.408756	0	0	1	0
2	1003	15	0	2	3	4.210570	26	0	2	0
3	1004	17	1	0	3	10.028829	14	0	3	1
4	1005	17	1	0	2	4.672495	17	1	3	0
5	1006	18	0	0	1	8.191219	0	0	1	1
6	1007	15	0	1	1	15.601680	10	0	3	0
7	1008	15	1	1	4	15.424496	22	1	1	1
8	1009	17	0	0	0	4.562008	1	0	2	0
9	1010	16	1	0	1	18.444466	0	0	3	1

```
In [ ]: # Display Last 10 rows
```

```
In [20]: data.tail(10)
```

```
Out[20]:
```

	StudentID	Age	Gender	Ethnicity	ParentalEducation	StudyTimeWeekly	Absences	Tutoring	ParentalSupport	Extracurricular
2382	3383	16	0	0	3	13.941823	20	0	2	
2383	3384	16	1	2	2	11.736409	18	1	4	
2384	3385	15	1	0	1	16.655581	13	1	3	
2385	3386	16	1	0	1	1.445434	20	0	3	
2386	3387	16	0	0	2	13.814021	14	0	2	
2387	3388	18	1	0	3	10.680555	2	0	4	
2388	3389	17	0	0	1	7.583217	4	1	4	
2389	3390	16	1	0	2	6.805500	20	0	2	
2390	3391	16	1	1	0	12.416653	17	0	2	
2391	3392	16	1	0	2	17.819907	13	0	2	

```
In [ ]: # Finding the Total no. of rows and columns
```

```
In [21]: data.shape
```

```
Out[21]: (2392, 15)
```

Data Preprocessing

```
In [ ]: # Checking null values
```

```
In [22]: data.isnull().sum()
```

```
Out[22]: StudentID      0
         Age            0
         Gender         0
         Ethnicity      0
         ParentalEducation 0
         StudyTimeWeekly 0
         Absences       0
         Tutoring       0
         ParentalSupport 0
         Extracurricular 0
         Sports         0
         Music          0
         Volunteering   0
         GPA            0
         GradeClass     0
         dtype: int64
```

```
In [ ]: # Removing empty rows
```

```
In [23]: data.dropna(inplace=True)
```

```
In [24]: data.shape # Size of rows and columns after removing empty rows .
```

```
Out[24]: (2392, 15)
```

```
In [25]: ### Feature Scaling
```

```
In [26]: numerical_features = ['Age', 'StudyTimeWeekly', 'Absences', 'GPA']
```

```
scaler = StandardScaler()
```

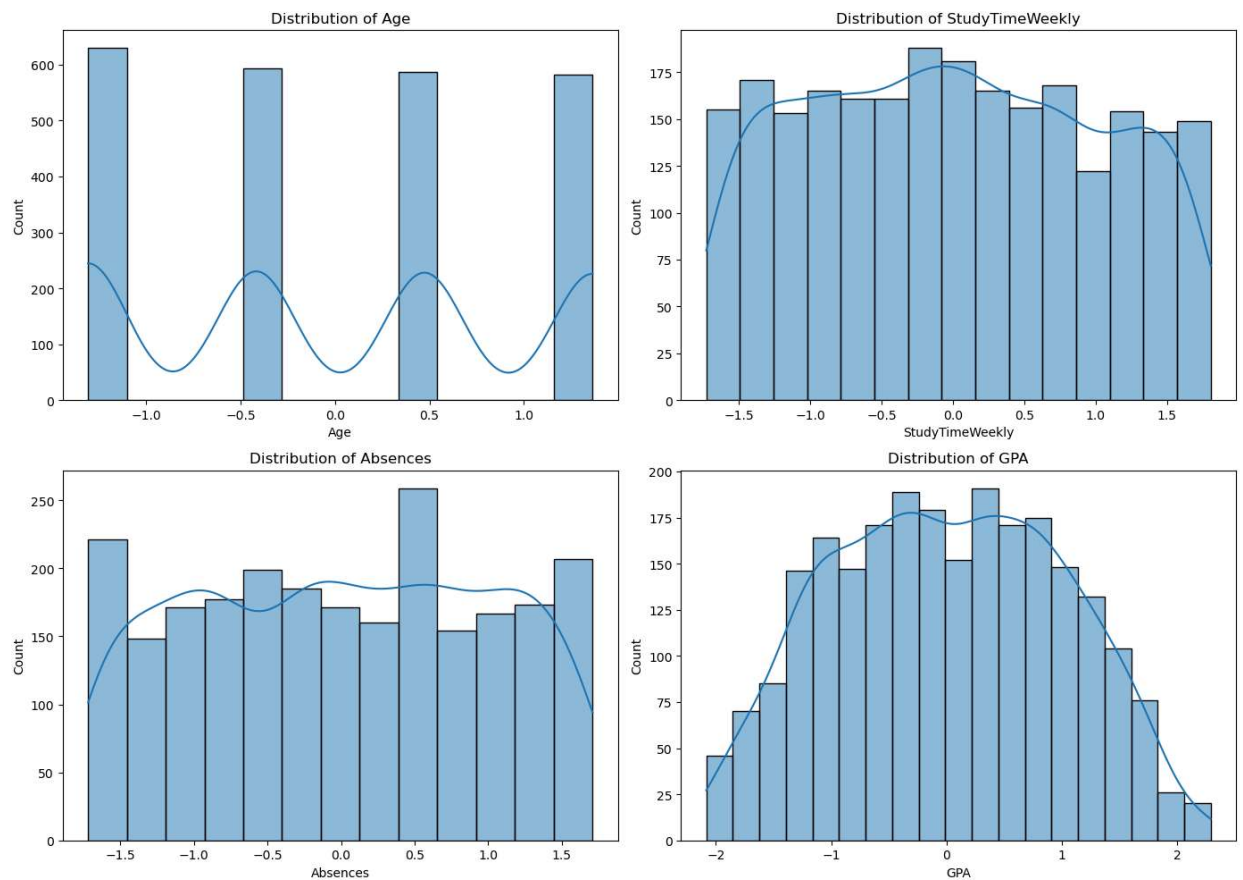
```
In [27]: data[numerical_features] = scaler.fit_transform(data[numerical_features])
         data.head(10)
```

```
Out[27]:
```

	StudentID	Age	Gender	Ethnicity	ParentalEducation	StudyTimeWeekly	Absences	Tutoring	ParentalSupport	Extracurr
0	1001	0.472919	1	0	2	1.780336	-0.890822	1		2
1	1002	1.362944	0	0	1	0.997376	-1.717694	0		1
2	1003	-1.307132	0	2	3	-0.984045	1.353542	0		2
3	1004	0.472919	1	0	3	0.045445	-0.063951	0		3
4	1005	0.472919	1	0	2	-0.902311	0.290422	1		3
5	1006	1.362944	0	0	1	-0.279704	-1.717694	0		1
6	1007	-1.307132	0	1	1	1.031513	-0.536449	0		3
7	1008	-1.307132	1	1	4	1.000161	0.881045	1		1
8	1009	0.472919	0	0	0	-0.921861	-1.599569	0		2
9	1010	-0.417106	1	0	1	1.534519	-1.717694	0		3


Exploratory Data Analysis (EDA)

```
In [28]: plt.figure(figsize=(14,10))
for i,feature in enumerate(numerical_features,1):
    plt.subplot(2,2,i)
    sns.histplot(data[data[feature]],kde=True)
    plt.title(f'Distribution of {feature}')
plt.tight_layout()
plt.show()
```



```
In [29]: ### Correlation Matrix
```

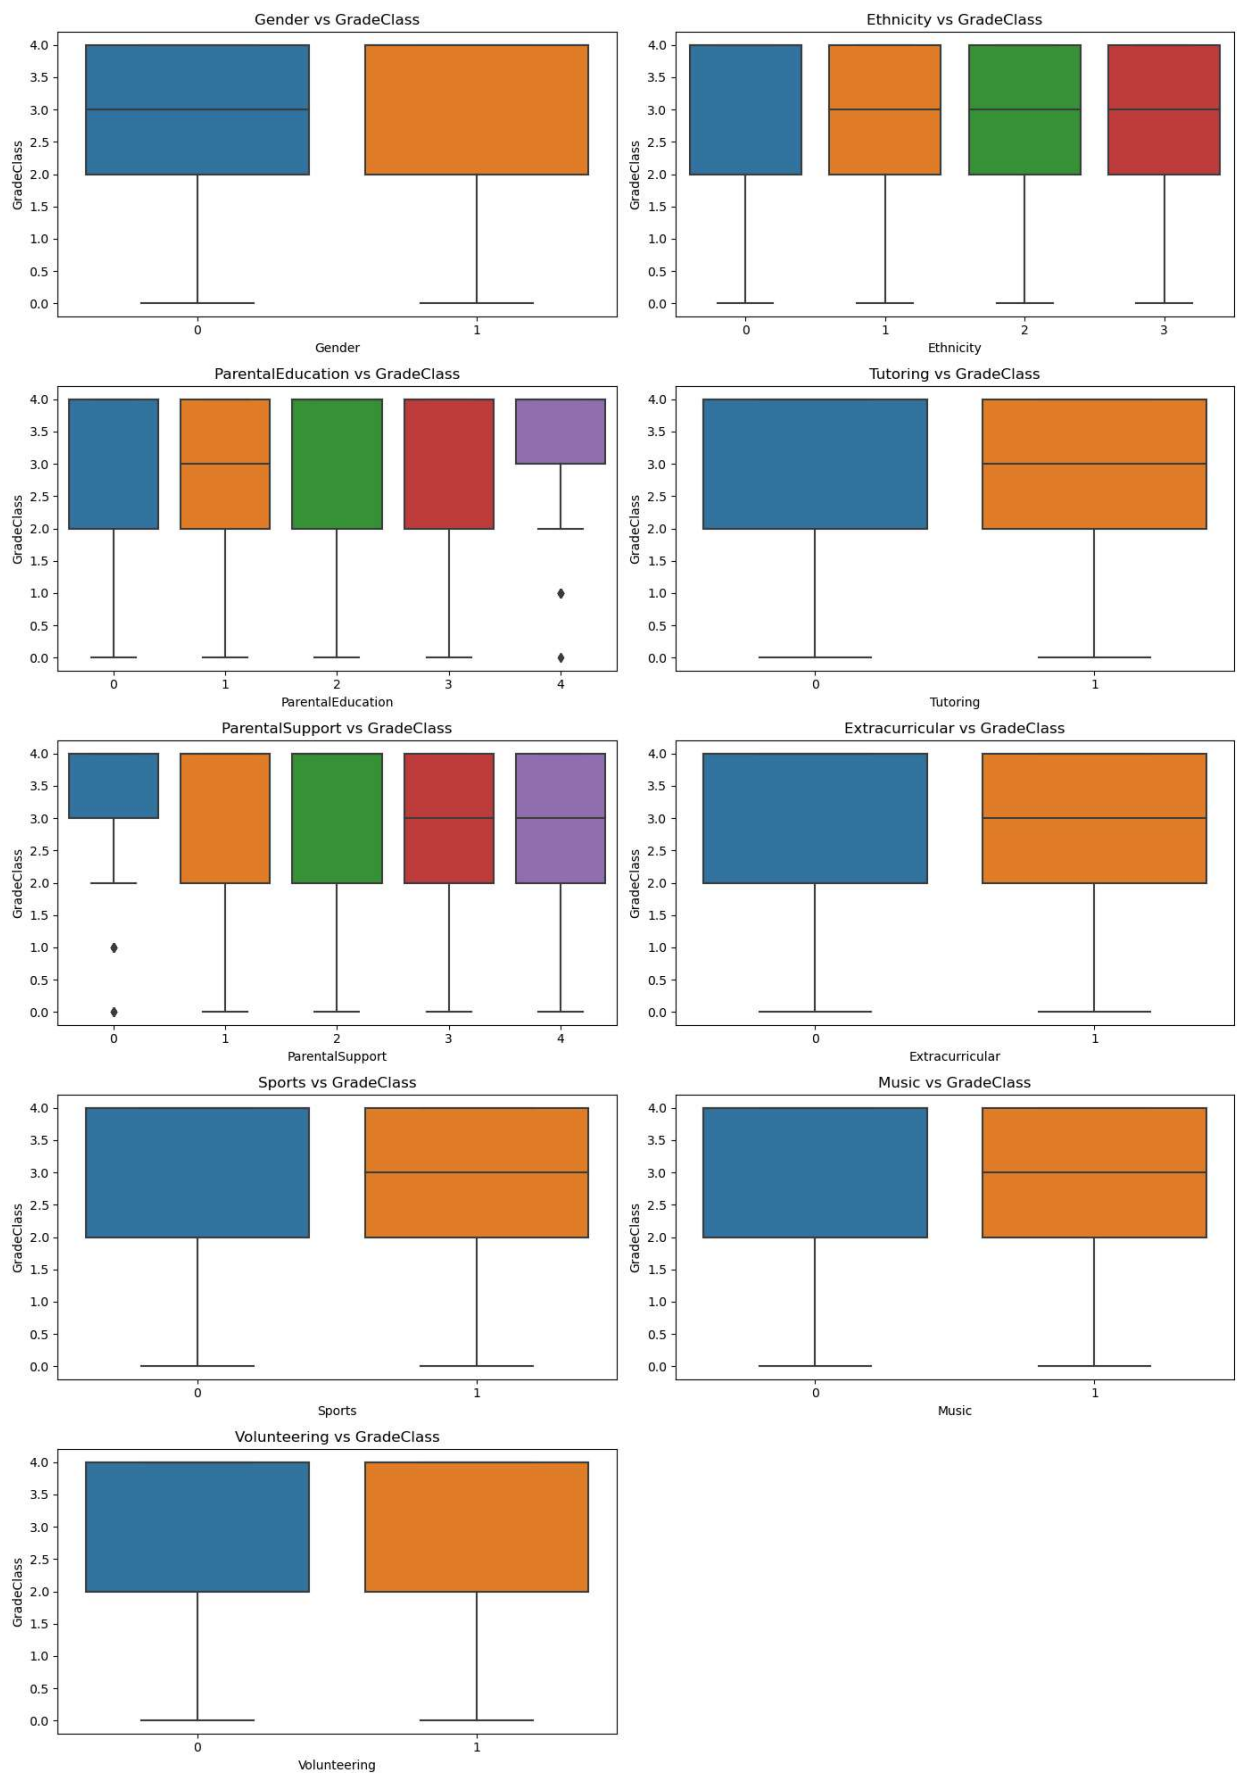
```
In [30]: plt.figure(figsize=(12,8))
corr_matrix = data.corr()
sns.heatmap(corr_matrix,annot = True,cmap='coolwarm',linewidths=0.5)
plt.title('Correlation Matrix')
plt.show()
```



```
In [31]: ### Box plots for Categorical Features
```

```
In [33]: categorical_features = ['Gender', 'Ethnicity', 'ParentalEducation', 'Tutoring', 'ParentalSupport', 'ExtracurricularActivities']

plt.figure(figsize=(14,20))
for i, feature in enumerate(categorical_features, 1):
    plt.subplot(5, 2, i)
    sns.boxplot(x=data[feature], y=data['GradeClass'])
    plt.title(f'{feature} vs GradeClass')
plt.tight_layout()
plt.show()
```



Model Selection and Training

In [34]: `### Split the data`

```
In [35]: x=data.drop(['StudentID','GradeClass'],axis=1)
y=data['GradeClass']
```

```
In [36]: ### Splitting the data into training and testing sets
```

```
In [37]: x_train, x_test, y_train, y_test = train_test_split(x,y,test_size=0.2,random_state=42)
```

Random Forest Classifier

```
In [50]: from sklearn.ensemble import RandomForestClassifier
RFC = RandomForestClassifier(random_state=42)
RFC.fit(x_train,y_train)
```

```
Out[50]: RandomForestClassifier
RandomForestClassifier(random_state=42)
```

Logistic Regression

```
In [49]: from sklearn.linear_model import LogisticRegression
LR = LogisticRegression(random_state=42,max_iter=1000) # max_iter avoids the warning
LR.fit(x_train,y_train)
```

```
Out[49]: LogisticRegression
LogisticRegression(max_iter=1000, random_state=42)
```

Support Vector Machine (SVM)

```
In [53]: from sklearn.svm import SVC as SVM
SVM = SVM(random_state = 42)
SVM.fit(x_train, y_train)
```

```
Out[53]: SVC
SVC(random_state=42)
```

K-Nearest Neighbour (KNN)

```
In [54]: from sklearn.neighbors import KNeighborsClassifier as KNN

kNN = KNN()
kNN.fit(x_train, y_train)
```

```
Out[54]: KNeighborsClassifier
KNeighborsClassifier()
```

Decision Tree Classifier

```
In [58]: from sklearn.tree import DecisionTreeClassifier

DTC = DecisionTreeClassifier(random_state=42)
DTC.fit(x_train, y_train)
```

```
Out[58]: DecisionTreeClassifier
DecisionTreeClassifier()
```

Model Evaluation

Evaluate the Model

i. Random Forest Classifier

```
In [75]: # Make predictions
y_pred = RFC.predict(x_test)

# Accuracy
accuracy = accuracy_score(y_test, y_pred)
print(f'Accuracy: {accuracy:.2f}')


# Classification Report
print("Classification Report")
print(classification_report(y_test, y_pred))

# Confusion Matrix
cm = confusion_matrix(y_test,y_pred)
plt.figure(figsize=(10,6))
sns.heatmap(cm,annot=True, fmt='d', cmap='Greys')
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title("Confusion Matrix")
plt.show()
```

Accuracy: 0.91

Classification Report

	precision	recall	f1-score	support
0.0	0.83	0.45	0.59	22
1.0	0.81	0.88	0.84	49
2.0	0.94	0.87	0.90	85
3.0	0.89	0.90	0.89	86
4.0	0.94	0.98	0.96	237
accuracy			0.91	479
macro avg	0.88	0.82	0.84	479
weighted avg	0.91	0.91	0.91	479



In []:

ii. Logistic Regression


```

In [74]: # Make prediction
y_pred=LR.predict(x_test)

# Accuracy
accuracy = accuracy_score(y_test,y_pred)
print(f'Logistic Regression Accuracy: {accuracy:.2f}')

# Classification Report
print("Logistic Regression Classification Report")
print(classification_report(y_test,y_pred))

# Confusion Matrix
cm = confusion_matrix(y_test,y_pred)
plt.figure(figsize=(10,6))
sns.heatmap(cm,annot=True, fmt='d', cmap='Blues')
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title("Logistic Regression Confusion Matrix")
plt.show()


```

```

Logistic Regression Accuracy: 0.75
Logistic Regression Classification Report

```

	precision	recall	f1-score	support
0.0	0.00	0.00	0.00	22
1.0	0.43	0.47	0.45	49
2.0	0.60	0.62	0.61	85
3.0	0.71	0.57	0.63	86
4.0	0.88	0.98	0.93	237
accuracy			0.75	479
macro avg	0.53	0.53	0.52	479
weighted avg	0.71	0.75	0.73	479



In []:

iii. Support Vector Machine (SVM)

```
In [81]: # Make prediction
y_pred=SVM.predict(x_test)

# Accuracy
accuracy = accuracy_score(y_test,y_pred)
print(f'SVM Accuracy: {accuracy:.2f}')
```

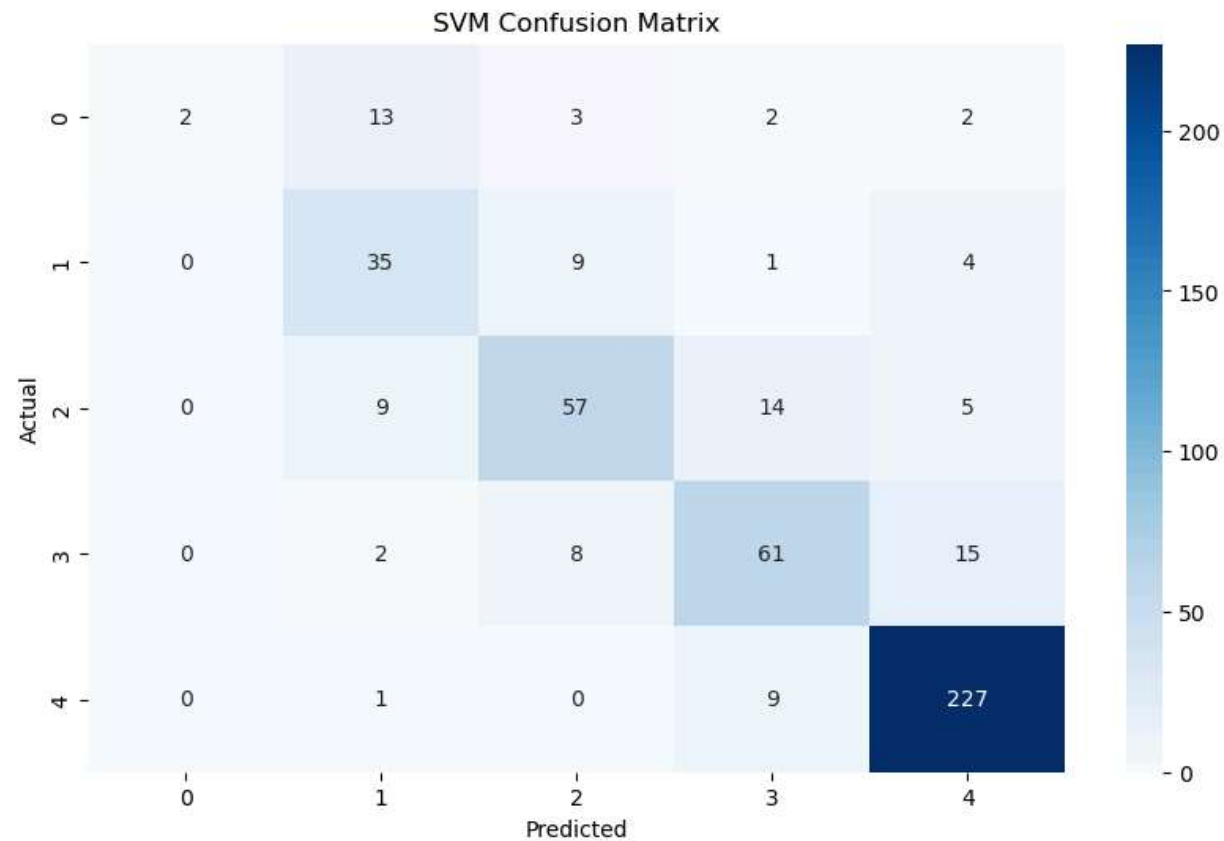
```
# Classification Report
print("SVM Classification Report")
print(classification_report(y_test,y_pred))
```

```
# Confusion Matrix
cm = confusion_matrix(y_test,y_pred)
plt.figure(figsize=(10,6))
sns.heatmap(cm,annot=True, fmt='d', cmap='Blues')
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title("SVM Confusion Matrix")
plt.show()
```

SVM Accuracy: 0.80

SVM Classification Report

	precision	recall	f1-score	support
0.0	1.00	0.09	0.17	22
1.0	0.58	0.71	0.64	49
2.0	0.74	0.67	0.70	85
3.0	0.70	0.71	0.71	86
4.0	0.90	0.96	0.93	237
accuracy			0.80	479
macro avg	0.78	0.63	0.63	479
weighted avg	0.81	0.80	0.78	479



```
In [ ]:
```

iv. k - Nearest Neighbors

```
In [66]: # Make prediction
y_pred=kNN.predict(x_test)

# Accuracy
accuracy = accuracy_score(y_test,y_pred)
print(f'kNN Accuracy: {accuracy:.2f}')
```

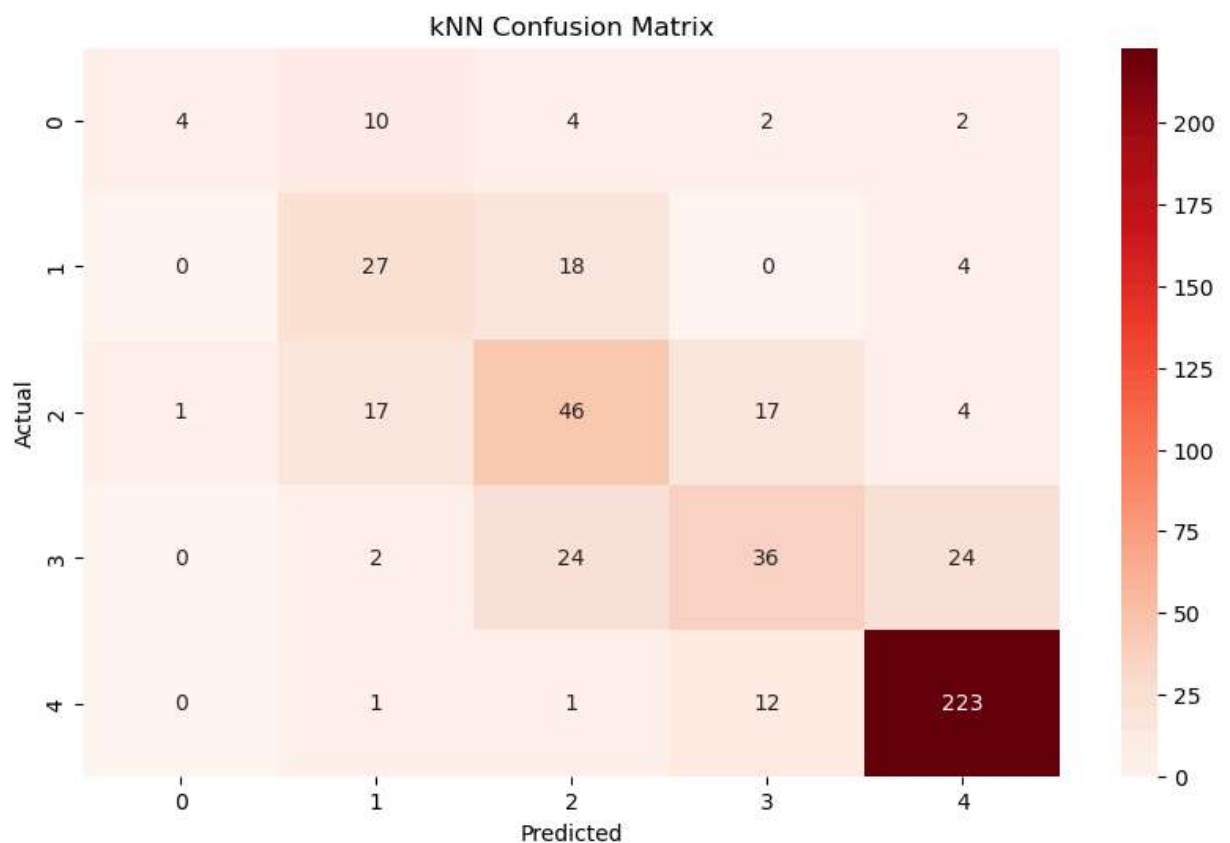
```
# Classification Report
print("kNN Classification Report")
print(classification_report(y_test,y_pred))
```

```
# Confusion Matrix
cm = confusion_matrix(y_test,y_pred)
plt.figure(figsize=(10,6))
sns.heatmap(cm,annot=True, fmt='d', cmap='Reds')
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title("kNN Confusion Matrix")
plt.show()
```

kNN Accuracy: 0.70

kNN Classification Report

	precision	recall	f1-score	support
0.0	0.80	0.18	0.30	22
1.0	0.47	0.55	0.51	49
2.0	0.49	0.54	0.52	85
3.0	0.54	0.42	0.47	86
4.0	0.87	0.94	0.90	237
accuracy			0.70	479
macro avg	0.63	0.53	0.54	479
weighted avg	0.70	0.70	0.69	479



In []:

v. Decision Tree Classifier

```
In [67]: # Make prediction
y_pred=DTC.predict(x_test)

# Accuracy
accuracy = accuracy_score(y_test,y_pred)
print(f'Decision Tree Accuracy: {accuracy:.2f}')


# Classification Report
print("Decision Tree Classification Report")
print(classification_report(y_test,y_pred))

# Confusion Matrix
cm = confusion_matrix(y_test,y_pred)
plt.figure(figsize=(10,6))
sns.heatmap(cm,annot=True, fmt='d', cmap='Greens')
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title("Decisioin Tree Confusion Matrix")
plt.show()
```

Decision Tree Accuracy: 0.84

Decision Tree Classification Report

	precision	recall	f1-score	support
0.0	0.60	0.55	0.57	22
1.0	0.66	0.80	0.72	49
2.0	0.86	0.75	0.81	85
3.0	0.81	0.80	0.81	86
4.0	0.90	0.91	0.90	237
accuracy			0.84	479
macro avg	0.77	0.76	0.76	479
weighted avg	0.84	0.84	0.84	479



In []: